

# MY472 - Data for Data Scientists

## Week 4: Textual Data

Thomas Robinson  
AT 2023

# Introduction

- This week we will focus on processing textual data
- Most file formats we work with in this course (.csv, .xml, .json, etc.) use text to store data
- The quantitative analysis of textual data is highly relevant in social science research and beyond
- We will discuss some basic analyses, but for a full course see [MY459](#) in Winter Term

# Plan for today

- Character encoding
- Text search: Globs and regular expressions
- Elementary text analysis
- Coding

# Character encoding

# Revisited: Basic units of data

- Bits
  - Smallest unit of storage; a 0 or 1
  - With  $n$  bits, can store  $2^n$  patterns
- Bytes
  - “eight bit encoding” represents characters through 8 bit, e.g. A represented as 65 = 01000001
  - 8 bits = 1 byte
  - Hence, 1 byte can store 256 patterns

# Encoding

- A “character set” is a list of characters with associated numerical representations
- The unique numbers associated with characters are called “code points”
- ASCII: The original character set, uses just 7 bits ( $2^7$ ), see <https://en.wikipedia.org/wiki/ASCII>
- ASCII was later extended, e.g. [ISO-8859](#), using 8 bits ( $2^8$ )
- Unfortunately, encoding has become a mess of differing standards, see [http://en.wikipedia.org/wiki/Character\\_encoding](http://en.wikipedia.org/wiki/Character_encoding)

# ASCII

Dec	Hx	Oct	Char	Dec	Hx	Oct	Html	Chr	Dec	Hx	Oct	Html	Chr	Dec	Hx	Oct	Html	Chr
0	0	000	<b>NUL</b> (null)	32	20	040	&#32;	<b>Space</b>	64	40	100	&#64;	<b>@</b>	96	60	140	&#96;	<b>`</b>
1	1	001	<b>SOH</b> (start of heading)	33	21	041	&#33;	<b>!</b>	65	41	101	&#65;	<b>A</b>	97	61	141	&#97;	<b>a</b>
2	2	002	<b>STX</b> (start of text)	34	22	042	&#34;	<b>"</b>	66	42	102	&#66;	<b>B</b>	98	62	142	&#98;	<b>b</b>
3	3	003	<b>ETX</b> (end of text)	35	23	043	&#35;	<b>#</b>	67	43	103	&#67;	<b>C</b>	99	63	143	&#99;	<b>c</b>
4	4	004	<b>EOT</b> (end of transmission)	36	24	044	&#36;	<b>\$</b>	68	44	104	&#68;	<b>D</b>	100	64	144	&#100;	<b>d</b>
5	5	005	<b>ENQ</b> (enquiry)	37	25	045	&#37;	<b>%</b>	69	45	105	&#69;	<b>E</b>	101	65	145	&#101;	<b>e</b>
6	6	006	<b>ACK</b> (acknowledge)	38	26	046	&#38;	<b>&amp;</b>	70	46	106	&#70;	<b>F</b>	102	66	146	&#102;	<b>f</b>
7	7	007	<b>BEL</b> (bell)	39	27	047	&#39;	<b>'</b>	71	47	107	&#71;	<b>G</b>	103	67	147	&#103;	<b>g</b>
8	8	010	<b>BS</b> (backspace)	40	28	050	&#40;	<b>(</b>	72	48	110	&#72;	<b>H</b>	104	68	150	&#104;	<b>h</b>
9	9	011	<b>TAB</b> (horizontal tab)	41	29	051	&#41;	<b>)</b>	73	49	111	&#73;	<b>I</b>	105	69	151	&#105;	<b>i</b>
10	A	012	<b>LF</b> (NL line feed, new line)	42	2A	052	&#42;	<b>*</b>	74	4A	112	&#74;	<b>J</b>	106	6A	152	&#106;	<b>j</b>
11	B	013	<b>VT</b> (vertical tab)	43	2B	053	&#43;	<b>+</b>	75	4B	113	&#75;	<b>K</b>	107	6B	153	&#107;	<b>k</b>
12	C	014	<b>FF</b> (NP form feed, new page)	44	2C	054	&#44;	<b>,</b>	76	4C	114	&#76;	<b>L</b>	108	6C	154	&#108;	<b>l</b>
13	D	015	<b>CR</b> (carriage return)	45	2D	055	&#45;	<b>-</b>	77	4D	115	&#77;	<b>M</b>	109	6D	155	&#109;	<b>m</b>
14	E	016	<b>SO</b> (shift out)	46	2E	056	&#46;	<b>.</b>	78	4E	116	&#78;	<b>N</b>	110	6E	156	&#110;	<b>n</b>
15	F	017	<b>SI</b> (shift in)	47	2F	057	&#47;	<b>/</b>	79	4F	117	&#79;	<b>O</b>	111	6F	157	&#111;	<b>o</b>
16	10	020	<b>DLE</b> (data link escape)	48	30	060	&#48;	<b>0</b>	80	50	120	&#80;	<b>P</b>	112	70	160	&#112;	<b>p</b>
17	11	021	<b>DC1</b> (device control 1)	49	31	061	&#49;	<b>1</b>	81	51	121	&#81;	<b>Q</b>	113	71	161	&#113;	<b>q</b>
18	12	022	<b>DC2</b> (device control 2)	50	32	062	&#50;	<b>2</b>	82	52	122	&#82;	<b>R</b>	114	72	162	&#114;	<b>r</b>
19	13	023	<b>DC3</b> (device control 3)	51	33	063	&#51;	<b>3</b>	83	53	123	&#83;	<b>S</b>	115	73	163	&#115;	<b>s</b>
20	14	024	<b>DC4</b> (device control 4)	52	34	064	&#52;	<b>4</b>	84	54	124	&#84;	<b>T</b>	116	74	164	&#116;	<b>t</b>
21	15	025	<b>NAK</b> (negative acknowledge)	53	35	065	&#53;	<b>5</b>	85	55	125	&#85;	<b>U</b>	117	75	165	&#117;	<b>u</b>
22	16	026	<b>SYN</b> (synchronous idle)	54	36	066	&#54;	<b>6</b>	86	56	126	&#86;	<b>V</b>	118	76	166	&#118;	<b>v</b>
23	17	027	<b>ETB</b> (end of trans. block)	55	37	067	&#55;	<b>7</b>	87	57	127	&#87;	<b>W</b>	119	77	167	&#119;	<b>w</b>
24	18	030	<b>CAN</b> (cancel)	56	38	070	&#56;	<b>8</b>	88	58	130	&#88;	<b>X</b>	120	78	170	&#120;	<b>x</b>
25	19	031	<b>EM</b> (end of medium)	57	39	071	&#57;	<b>9</b>	89	59	131	&#89;	<b>Y</b>	121	79	171	&#121;	<b>y</b>
26	1A	032	<b>SUB</b> (substitute)	58	3A	072	&#58;	<b>:</b>	90	5A	132	&#90;	<b>Z</b>	122	7A	172	&#122;	<b>z</b>
27	1B	033	<b>ESC</b> (escape)	59	3B	073	&#59;	<b>;</b>	91	5B	133	&#91;	<b>[</b>	123	7B	173	&#123;	<b>{</b>
28	1C	034	<b>FS</b> (file separator)	60	3C	074	&#60;	<b>&lt;</b>	92	5C	134	&#92;	<b>\</b>	124	7C	174	&#124;	<b> </b>
29	1D	035	<b>GS</b> (group separator)	61	3D	075	&#61;	<b>=</b>	93	5D	135	&#93;	<b>]</b>	125	7D	175	&#125;	<b>}</b>
30	1E	036	<b>RS</b> (record separator)	62	3E	076	&#62;	<b>&gt;</b>	94	5E	136	&#94;	<b>^</b>	126	7E	176	&#126;	<b>~</b>
31	1F	037	<b>US</b> (unit separator)	63	3F	077	&#63;	<b>?</b>	95	5F	137	&#95;	<b>_</b>	127	7F	177	&#127;	<b>DEL</b>

Source: [www.LookupTables.com](http://www.LookupTables.com)

# Potential encoding issues

(Wrongly) detected encoding:

- Encoding type/character set is not stored as metadata in plain text files
- So software guesses which encoding is used, which might go wrong
- Assuming the wrong encoding when reading in/parsing a text file leads to import errors and corrupted characters ([Mojibake](#)): Underlying bit sequences are translated into the wrong characters

Space:

- 8 bits are too few to store all known characters
- Encoding with 32 bits, however, would imply a lot of rarely used bits
- Those bits take up memory, implying unnecessarily large file sizes



# Widely used character encoding today: Unicode

- Created by the [Unicode Consortium](#)
- Common Unicode encoding formats: **UTF-8** and **UTF-16** (Unicode transformation format)
- UTF-8 is a variable-width character encoding and by far the most frequent character encoding on the web today
- Variable amounts of bits are used for each character with the first byte/8 bits corresponding to ASCII
- Common characters therefore need less space, but system capable of storing vast amounts of character code points

# UTF-8 details

Number of bytes	Byte 1	Byte 2	Byte 3	Byte 4
1	0xxxxxxx			
2	110xxxxx	10xxxxxx		
3	1110xxxx	10xxxxxx	10xxxxxx	
4	11110xxx	10xxxxxx	10xxxxxx	10xxxxxx

---

<https://en.wikipedia.org/wiki/UTF-8>

Try it out: Create two .txt files, one containing a single line with the character *ä*, the other one a single line with the character *ü*. Then check the sizes of both files in bytes which should be different if files are encoded in UTF-8.

# Things to watch out for

- Many text production softwares (e.g. MS Office-based products) might still use proprietary character encoding formats, such as Windows-1252
- Windows tends to use UTF-16, while Unix-based platforms use UTF-8
- Text editors can be misleading: the client may display mojibake but the encoding might still be as intended
- Generally, no easy method of detecting encodings in basic text files

# Some things to try with encoding issues

To determine the estimated character encoding of a file (note that this estimate might be incorrect)

- Linux, Unix, Mac: For example, `file -I filename.txt`, `file -I filename.json`, etc. in terminal
- Windows: For example, open with Notepad and check field in the lower right hand corner of the window

To change a file's encoding (see e.g. this [Stack Overflow post](#))

- Linux, Unix, Mac: For example, `iconv -f ISO-8859-15 -t UTF-8 in.txt > out.txt` in terminal
- Windows: For example, open the text with Notepad, click "Save As", and choose a name and UTF-8 encoding. Alternatively, use PowerShell

# Some things to try with encoding issues (in R)

In R, e.g. via `readr` (for more discussion, see [R4DS](#))

- For a character vector `x`, obtain texts assuming a different encoding with `parse_character(x, locale = locale(encoding = "Latin1"))`
- Make guess about encoding with `guess_encoding(charToRaw(x))`

# Globs and regular expressions

# Globs

- Searching and counting specific words in texts is key for quantitative textual analysis
- Globbs offer a simple and intuitive approach to search through text with wildcard characters
- Glob patterns originally used to search file and folder names

# Globs: Exemplary syntax

Wildcard	Description	Examples	Exemplary matches
*	Any number (also zero) of characters	tax*, *tax*	taxation, overtaxed
?	Single character	??flation	inflation or deflation
[ab], [AB], [17], etc.	List of characters	module-[17].Rmd	module-1.Rmd or module-7.Rmd
[a-z], [A-Z], [0-9]	Range of characters	module-[A-Z].Rmd	module-A.Rmd or module-B.Rmd or module-C.Rmd ...

---

[https://en.wikipedia.org/wiki/Glob\\_\(programming\)](https://en.wikipedia.org/wiki/Glob_(programming))



# Regular expressions

- Powerful and much more flexible tool to search (and replace) text
- Different syntax than globs
- Text editors (e.g. VS Code) can usually find and replace terms with regular expressions
- Can also be used in many programming languages, e.g. when counting or collecting certain keywords in text analysis
- In R, we can e.g. use `stringr` or `quanteda` to search for keywords with regular expressions
- Topic could fill lectures itself, we will cover some basics here

# Sample text

Inflation in the Eurozone

2pm

2:30pm

2.15pm

2 15

11.30

22-30

5-15pm

Münster

Muenster

Munster

@

@JoeBiden

@KamalaHarris

# Regular expressions: Syntax

- Regular expressions can consist of literal characters and metacharacters
- **Literal characters:** Usual text
- **Metacharacters:** `^ $ [] () {} * + . ?` etc.
- When a meta character shall be treated as usual text in a search, escape it with (unless it is in a set `[]`) `\`
- For example, searching `.` in regex notation will select any character, but searching `\.` will select the actual full stop character

# Syntax: Specifying characters (1/2)

- `.`: Matches any character (also white spaces)
- `\d`: Matches any digit 0-9
- `\w`: Matches any character a-z, A-Z, 0-9, `_`
- `\s`: Matches white spaces
- Capitalised versions negate: `\S` matches everything that is not a white space etc.

# Syntax: Specifying characters (2/2)

- `^`: Matches characters at the beginning of the line or string,
  - E.g. `^M` will select all capital m at the beginning of strings or lines
- `$`: Matches characters at the end of the line or string,
  - E.g. `m$` will select all lowercase m at the end of strings or lines
- `[]`: Character set, e.g. `[a-zA-Z]` selects single characters from the Latin alphabet in lower and upper case letters, `[ai]` selects characters that are “a” or “i”, `[0-9]` digits from 0 to 9
- `[^ ]`: In brackets, `^` has a different meaning namely “not”, e.g. `[^a-z]` selects all characters that are not from the lower case alphabet

# Syntax: Selecting sequences of characters

In order to select whole words, we need to add quantifiers to individual characters:

- \*: Zero or more times, e.g. `in[a-z]*` will select *in* and also *inflation* in a search;
  - We could use `.*` represents all characters and white spaces
- +: One or more times, e.g. `in[a-z]+` will not select *in* but *inflation*
- ?: Denotes optional characters, e.g. `re?ally` will select *really* and *rally*
- {}: Specifies lengths of sequences, e.g. `\d{3}` selects sequences of 3 digits, `\w{3,4}` selects sequences between 3 and 4 general characters, and `\d{3,}` selects sequences of at least 3 digits

# Syntax: Boolean or and capturing groups

- |: Boolean or
- (): Capturing groups, e.g. (ue? | ü) selects u, ue, and ü.
  - This means that when searching text, the regular expression M(ue? | ü)nster will find *Münster*, *Muenster*, and *Munster*.
  - The captured groups can also be referenced with integer counts, which can be very helpful when replacing text
- [https://en.wikipedia.org/wiki/Regular\\_expression](https://en.wikipedia.org/wiki/Regular_expression)

# Regular expressions in R and beyond

- Regular expressions are used for flexible word searches in the `quantda` package
- `stringr` is another good package for strings that uses regular expressions:
  - `str_view()` show results of searches with regular expressions
  - `str_extract()` allows you to extract keywords from strings through regular expressions
  - `str_replace()` finds and replaces regular expressions
- Detailed discussion of strings and regular expressions with `stringr` in R [here](#)
- R markdown with many examples [here](#)



# More resources

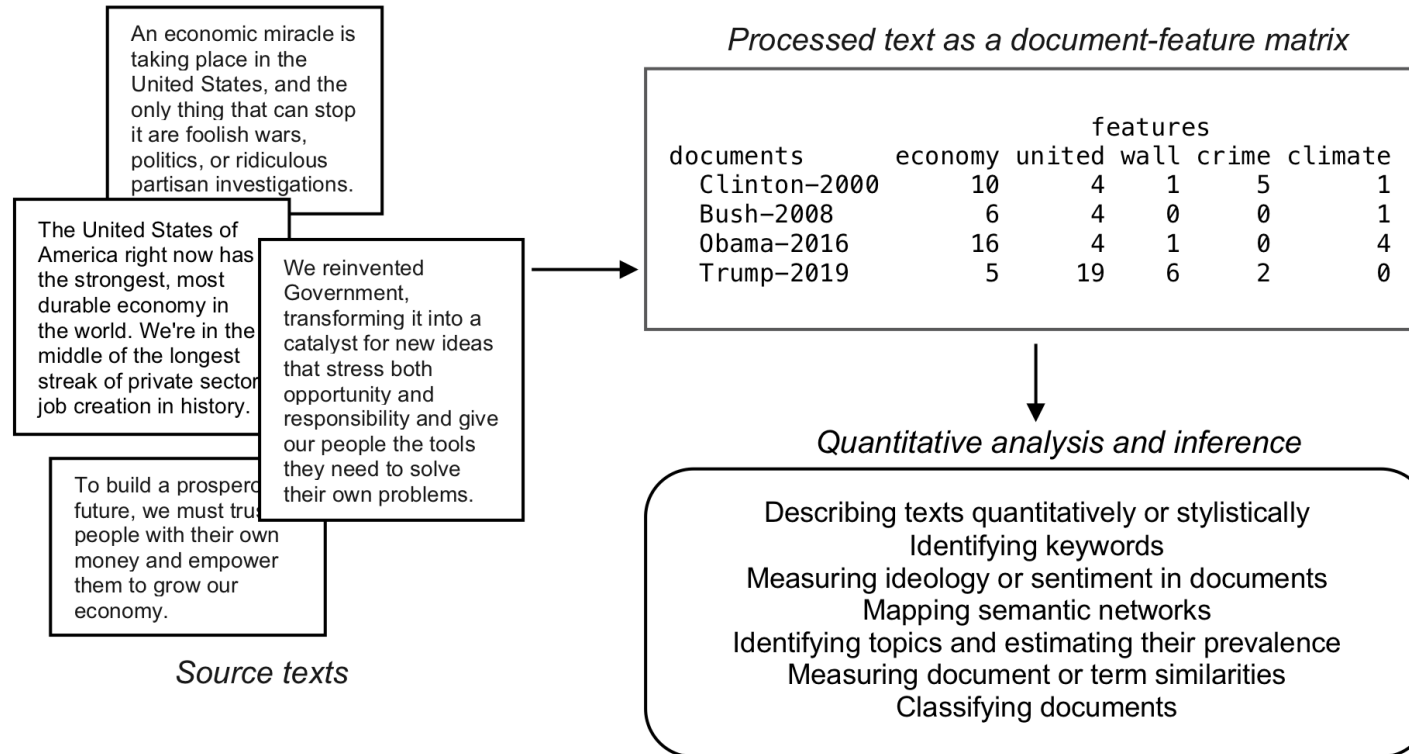
- Some good general discussions of the topic also on Youtube, e.g. [here](#)
- In depth treatment of regular expression (programming language independent): [\*Mastering Regular Expressions\*](#) by Jeffrey E. F. Fried
- A good place to test regular expressions and see the results visually is [regxr.com](#)
  - You can provide sample text, write a regex, and it will highlight matches

# Elementary text analysis

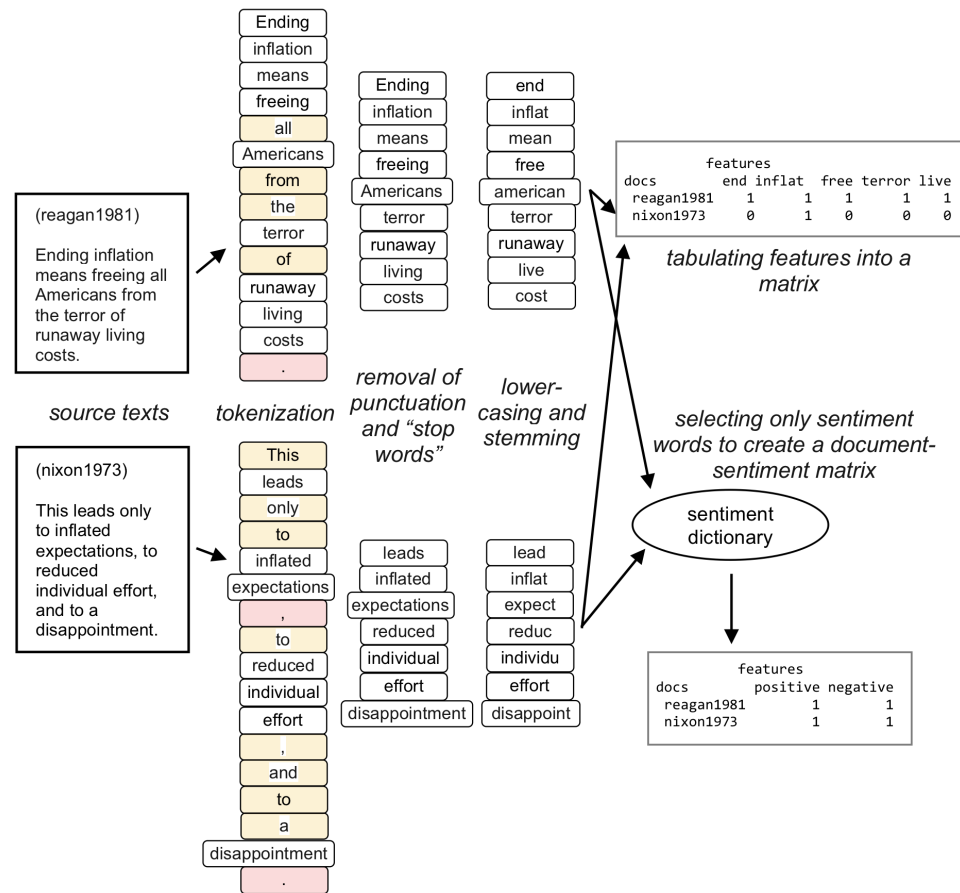
# Moving from texts to numbers

- To analyse text quantitatively, the key question is how to move from text to numbers
- We will look at very common approaches that count words in documents
- This abstracts from the sequential dependency of words (beyond n-grams) and is sometimes referred to as a bag-of-words approach

# Common workflow



# Common workflow: Tokenisation + dictionary method



# Some key concepts

- Document-feature matrix (dfm): As many rows as documents, as many columns as words/features after cleaning
- Stopwords: Common words such as “the”, “to”, etc.
- Stemming: Heuristic process to obtain the stem of words which in essence groups terms, see the following [link](#) for a detailed discussion
- n-grams: Sequences of words, e.g. bigrams (2) or trigrams (3). For example allows to record “not good” as a feature

# Dictionary approaches

- Map each word or phrase to a “dictionary” of words, e.g. associated with a known “sentiment” or psychological state or with certain topics
- Treats matches within each dictionary as equivalent
- Examples: Linguistic Inquiry and Word Count, or the General Inquirer

# Dictionary example (from LIWC 2015)

Dictionary object with 1 key entry.

- [posemo]:

- like, like\*, :), (:, accept, accepta\*, accepted, accepting, accepts, active, ...  
interests, invigor\*, joke\*, joking, jolly, joy\*, keen\*, kidding,  
kind, kindly, kindn\*, kiss\*, laidback, laugh\*, legit, libert\*,  
likeab\*, liked, likes, liking, livel\*, lmao\*, lmfaao\*, lol, love, loved, lovelier, ...



# Problems with dictionary approaches

- Polysemy – multiple meanings: The word “kind” has three!
- From State of the Union corpus: 318 matches
  - kind/NOUN – 95%
  - kind (of)/ADVERB – 1%
  - kind/ADJECTIVE – 4%
- These are known as false positives
- Other problem: False negatives (what we miss)
  - Missed: kindness
  - Also missed: altruistic and magnanimous
- How to treat conflicting keywords in the same string? “Had a great day ... not.”

# Further topics

- Text classification: How do we use a feature matrix to predict document labels (e.g. spam/not spam)?
- Topic models: How do we find sets of words which tend to appear together?
- Word and document embeddings: How can we represent words or documents as vectors and analyse their distances/similarities?
- How do we take into account the sequential nature of text?
- etc.

Coding

# Markdown files

- 01-regular-expressions-in-r.Rmd
- 02-text-analysis.Rmd
- 03-parsing-pdfs.Rmd