

MY472 - Data for Data Scientists

Week 10: NoSQL and Cloud Databases

Daniel de Kadt

28 November 2023

Outline

- Cloud solutions for databases
- SQL vs. noSQL
- Coding session
 - Online database example with SQL: BigQuery
 - NoSQL example: MongoDB

Cloud solutions

Why remote solutions?

- Last week we learned about relational databases
- Worked with SQL to manipulate data stored within tables
- In our applications, the data were **local**
- At scale, we invariably want to store data **remotely**
- Trade-offs, as always!

Some exemplary services

| Database Type | AWS | GCP | Azure |
|--------------------------|------------|---------------|--------------|
| Managed RDS | Amazon RDS | Cloud SQL | Azure SQL |
| Data Warehousing | Redshift | BigQuery | Snowflake |
| NoSQL (simple key-value) | DynamoDB | BigTable | Azure Tables |
| NoSQL (document) | DocumentDB | MongoDB on GC | Cosmos DB |

Google Cloud Platform: BigQuery

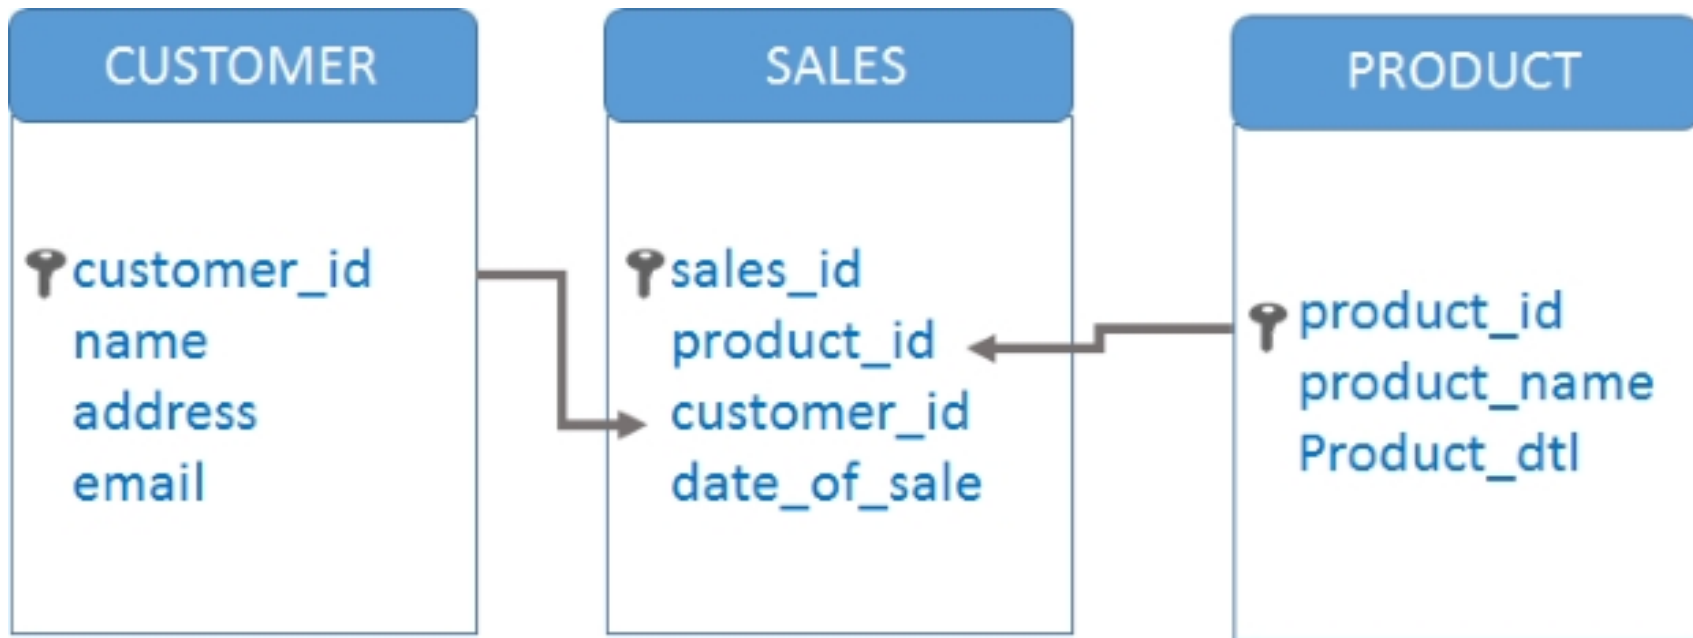
- To create and query online databases, we will look at Google BigQuery's sandbox version as an example
- Database warehouse with other features, used by many financial and commercial companies
- Queried via SQL syntax (API access allows integration with R or Python)
- Scalable to very large databases
- Good [documentation](#)
- Many similar databases exist from other providers

SQL vs noSQL

SQL

- Relational databases have a strict structure

A simple e-commerce example:

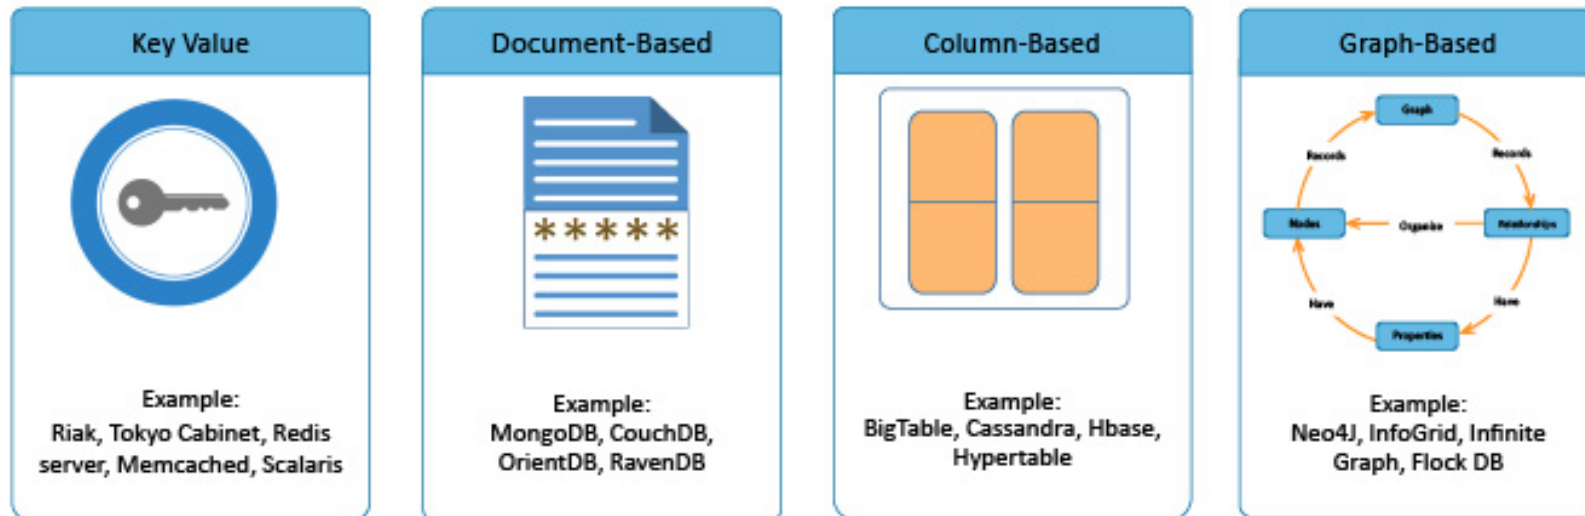


noSQL

- Originally referring to “non SQL”, “non relational” or “not only SQL”
- Provides a mechanism for storage and retrieval of data which is modeled in means other than the tabular relations used in relational databases
- No strict structure/schema
- noSQL databases are good for data with
 - High **velocity** – Lots of data coming in very quickly
 - High **variety** – Data can be structured, semi-structured, and unstructured
 - High **volume** – Total size of data
 - High **complexity** – Stored in many locations

noSQL types

Some examples from recent years:



© Simplelearn. All rights reserved.

simplelearn

From: [Simplelearn](https://www.simplelearn.com/)

noSQL: Pros and Cons

PROS

Massive scalability

High availability

Schema flexibility

Sparse and semistructured data

CONS

Limited query capabilities

Not standardized

Not matured

Developer heavy

MongoDB

- **Document-based** database
- Mapping of concepts
- Each document is constructed as a **BSON** ([Binary JSON](#))
- Not UTF-8 string encoded document
- Like JSON, but binary - machine readable only (very lightweight)
- Can store more data types: Dates, separate kinds of numerics (int, float, etc.)

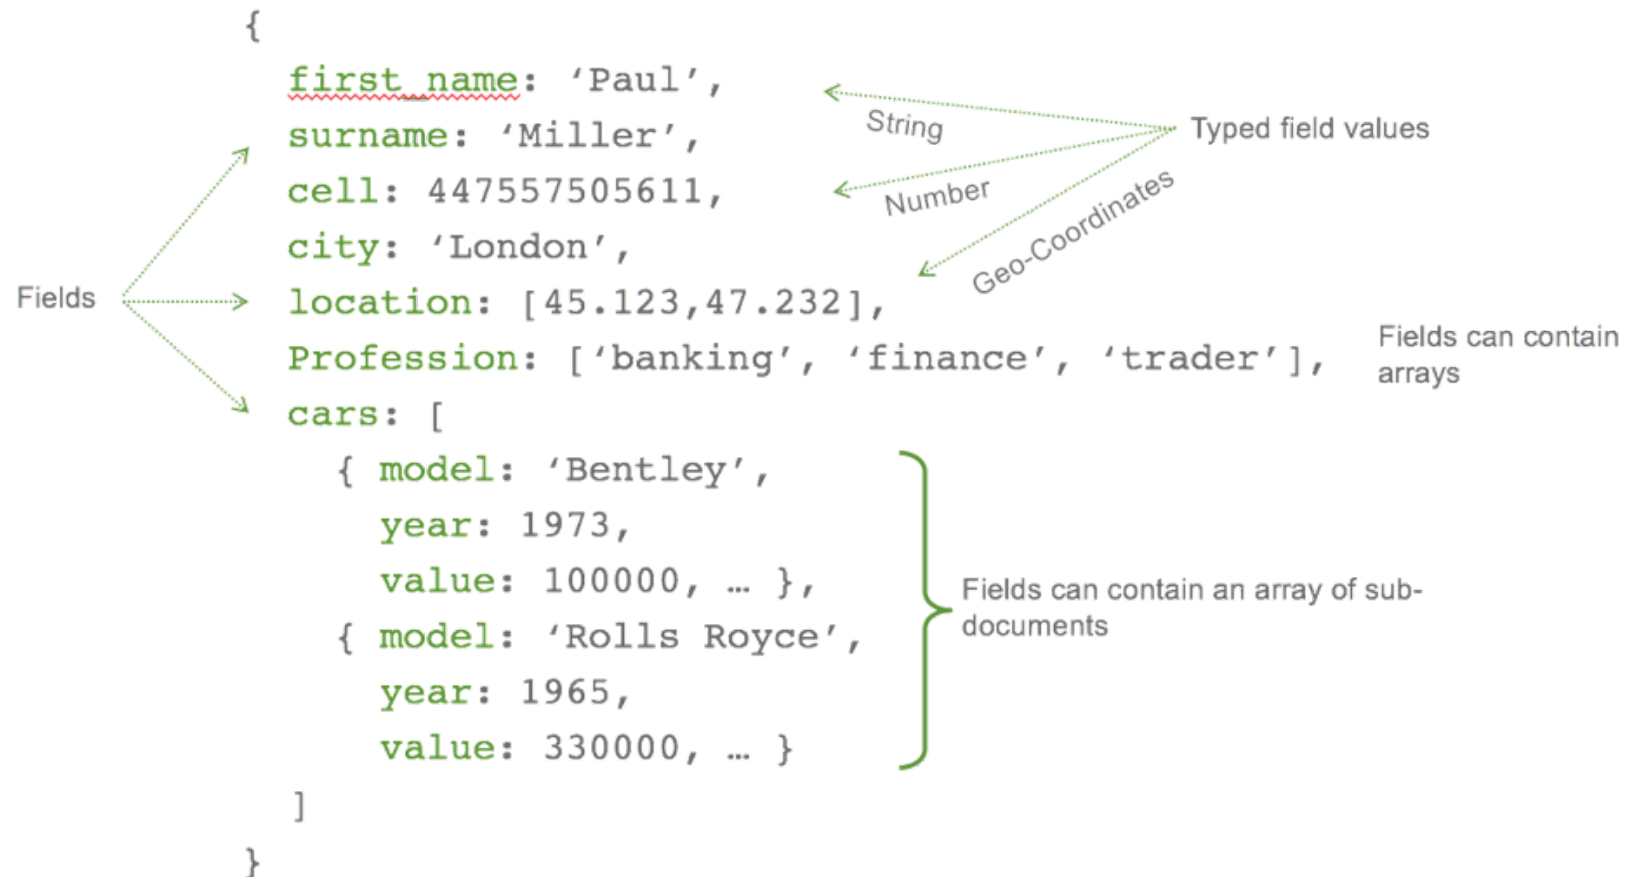
[Reference](#)

MongoDB vs. SQL

| SQL Terms/Concepts | MongoDB Terms/Concepts |
|---|---|
| database | database |
| table | collection |
| row | document or BSON document |
| column | field |
| index | index |
| table joins | <code>\$lookup</code> , embedded documents |
| primary key | primary key |
| Specify any unique column or column combination as primary key. | In MongoDB, the primary key is automatically set to the <code>_id</code> field. |

MongoDB documents

A document looks like this:



From: datawow.io

MongoDB in R (optional)

- We will look at MongoDB as an example of a popular noSQL database this week
- We thereby try to replicate basic queries from last week using MongoDB via R with the package `mongolite`
- For a simple selection of documents (i.e. rows in SQL), we will use its `find()` method
- For a bit more sophisticated queries, we will use the `aggregate()` method
- Search queries are in JSON like notation
- Detailed [documentation](#) of MongoDB commands and operators
- [Resource 1 \(pdf\)](#) and [resource 2 \(website\)](#) for the R package `mongolite`

Coding session

Files this week

- 01-bigquery-create-own-database.Rmd
- 02-bigquery-examples.Rmd
- 03-mongodb-demo.Rmd